

# Heuristic Algorithms

Master's Degree in Computer Science/Mathematics

Roberto Cordone

DI - Università degli Studi di Milano



Schedule: Thursday 14.30 - 16.30 on MS-Teams

Friday 14.30 - 16.30 on MS-Teams

Office hours: on appointment

E-mail: [roberto.cordone@unimi.it](mailto:roberto.cordone@unimi.it)

Web page: <https://homes.di.unimi.it/cordone/courses/2020-ae/2020-ae.html>

Ariel site: <https://rcordoneha.ariel.ctu.unimi.it>

# Overcoming local optima

The *steepest descent* exchange heuristics only provide local optima

In order to improve, one can

- repeat the search (*How to avoid following the same path?*)
- extend the search (*How to avoid falling in the same optimum?*)

In the constructive algorithms only repetition was possible

The constructive metaheuristics exploit

- randomization
- memory

to operate on  $\Delta_A^+(x)$  and  $\varphi_A(i, x)$

The exchange metaheuristics exploit them to operate on

- 1 the starting solution  $x^{(0)}$  (multi-start, *ILS*, *VNS*)
- 2 the neighbourhood  $N(x)$  (*VND*)
- 3 the selection criterium  $\varphi(x, A, D)$  (*DLS* or *GLS*, *SA*, *TS*)

# Termination condition

A search that repeats or proceeds beyond local optimum can ideally be infinite

In practice, one uses “absolute” termination conditions

- 1 a given total number of explorations of the neighbourhood or a given total number of repetitions of the local search
- 2 a given total execution time
- 3 a given value of the objective
- 4 a given improvement of the objective with respect to the starting solution

or “relative” termination conditions

- 1 a given number of explorations of the neighbourhood or repetitions after the last improvement of  $f^*$
- 2 a given execution time after the last improvement
- 3 a given minimum value of the ratio between improvement of the objective and number of explorations or execution time  
(e.g.:  $f^*$  improves less than 1% in the last 1000 explorations)

Fair comparisons require absolute conditions (time or number of explorations)

# Modify the starting solution: random generation

It is possible to create different starting solutions

- generating them at random
- applying different constructive heuristics
- modifying solutions generated by the exchange algorithm

The advantages of random generation are

- conceptual simplicity
- quickness for the problems in which it is easy to guarantee feasibility
- control on the probability distribution in  $X$  based on
  - element cost (e.g., favour the cheapest elements)
  - element frequency during the past search, to favour the most frequent elements (intensification) or the less frequent ones (diversification)

*This combines randomization and memory*

- asymptotic convergence to the optimum (in infinite time)

The disadvantages of random generation are

- scarce quality of the starting solutions (*not the final ones!*)
- long times before reaching the local optimum

*This depends on the complexity of the exchange algorithm*

- inefficiency when deciding feasibility is  $\mathcal{NP}$ -complete

# Modify the starting solution: constructive procedures

**Multi-start** methods are the classical approach

- design several constructive heuristics
- each constructive heuristic generates a starting solution
- each starting solution is improved by the exchange heuristic

The disadvantages are

- 1 **scarce control**: the generated solutions tend to be similar
- 2 **impossibility to proceed indefinitely**: the number of repetitions is fixed
- 3 **high design effort**: several different algorithms must be designed
- 4 **no guarantee of convergence**, not even in infinite time

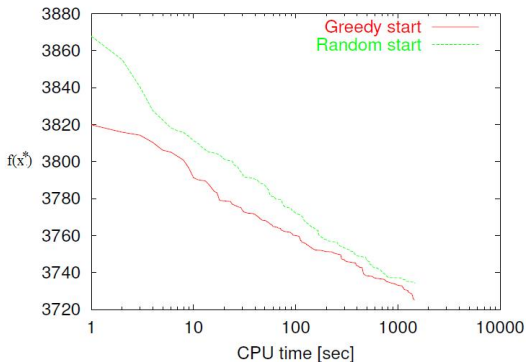
Consequently, constructive metaheuristics are preferred nowadays

*GRASP* and Ant System include by definition an exchange procedure

# Influence of the starting solution

If the exchange heuristic is

- **good**, the starting solution has a short-lived influence:  
a random or heuristic generation of  $x^{(0)}$  are very similar
- **bad**, the starting solution has a long-lived influence:  
a good heuristic to generate  $x^{(0)}$  is useful



*This exchange heuristic is not very good*

# Modify the starting solution exploiting the previous ones

The idea is to exploit the information on previously visited solutions

- save reference solutions, such as the best local optimum found so far and possibly other local optima
- generate the new starting solution modifying the reference ones

The advantages of this approach are

- control: the modification can be reduced or increased *ad libitum*
- good quality: the starting solution is very good
- conceptual simplicity
- implementation simplicity: the modification can be performed with the operations defining the neighbourhood
- asymptotic convergence to the optimum under suitable conditions

# Iterated Local Search (*ILS*)

The Iterated Local Search (*ILS*) requires

- an *steepest descent* exchange heuristic to produce local optima
- a **perturbation procedure** to generate the starting solutions
- an **acceptance condition** to decide whether to change the reference solution  $x$
- a termination condition

*Algorithm* IteratedLocalSearch( $I, x^{(0)}$ )

$x := \text{SteepestDescent}(x^{(0)}); x^* := x;$

*For*  $l := 1$  *to*  $\ell$  *do*

$x' := \text{Perturbate}(x);$

$x' := \text{SteepestDescent}(x');$

*If* **Accept**( $x', x^*$ ) *then*  $x := x';$

*If*  $f(x') < f(x^*)$  *then*  $x^* := x';$

*EndWhile*;

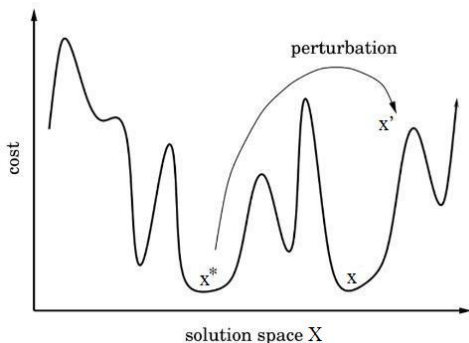
*Return* ( $x^*, f(x^*)$ );



# Iterated Local Search (ILS)

The idea is that

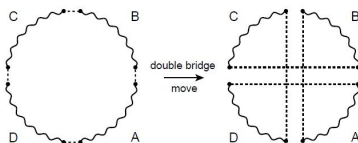
- the exchange heuristic quickly explores an attraction basin, terminating into a local optimum
- the perturbation procedure moves to another attraction basin
- the acceptance condition evaluates if the new local optimum is a promising starting point for the following perturbation



# Example: *ILS* for the *TSP*

A classical application of *ILS* to the *TSP* uses

- exchange heuristic: *steepest descent* with neighbourhood  $N_{\mathcal{R}_2}$  or  $N_{\mathcal{R}_3}$
- perturbation procedure: a *double-bridge* move that is particular kind of 4-exchange



- acceptance condition: the best known solution improves

$$f(x') < f(x^*)$$

# Perturbation procedure

Let  $\mathcal{O}$  be the operation set that defines neighbourhood  $N_{\mathcal{O}}$

The **perturbation procedure** performs a random operation  $o$

- with  $o \in \mathcal{O}' \not\subseteq \mathcal{O}$ , to avoid that the exchange heuristic drive solution  $x'$  back to the starting local optimum  $x$

Two typical definitions of  $\mathcal{O}'$  are

- sequences of  $k > 1$  operations of  $\mathcal{O}$   
(generating a random sequence is cheap)
- conceptually different operations  
(e.g., vertex exchanges instead of arc exchanges)

The main difficulty of *ILS* is in **tuning the perturbation**: if it is

- too strong, it turns the search into a random restart
- too weak, it guides the search back to the starting optimum
  - wasting time
  - possibly losing the asymptotic convergence

Ideally one would like to **enter any basin** and **get out of any basin**

# Acceptance condition

```
Algorithm IteratedLocalSearch( $I, x^{(0)}$ )  
 $x := \text{SteepestDescent}(x^{(0)}); x^* := x;$   
For  $l := 1$  to  $\ell$  do  
     $x' := \text{Perturbate}(x);$   
     $x' := \text{SteepestDescent}(x');$   
    If  $\text{Accept}(x', x^*)$  then  $x := x';$   
    If  $f(x') < f(x^*)$  then  $x^* := x';$   
EndWhile;  
Return  $(x^*, f(x^*));$ 
```

The acceptance condition balances intensification and diversification

- accepting only improving solutions favours intensification

$$\text{Accept}(x', x^*) := (f(x') < f(x^*))$$

The reference solution is always the best found:  $x = x^*$

- accepting any solution favours diversification

$$\text{Accept}(x', x^*) := \text{true}$$

The reference solution is always the last optimum found:  $x = x'$

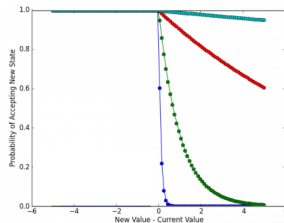
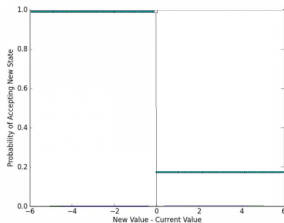
# Acceptance condition

Intermediate strategies can be defined based on  $\delta f = f(x') - f(x^*)$

- if  $\delta f < 0$ , always accept  $x'$
- if  $\delta f \geq 0$ , accept  $x'$  with probability  $\pi(\delta f)$ ,  
where  $\pi(\cdot)$  is a nonincreasing function

The most typical cases are:

- constant probability:  $\pi(\delta f) = \bar{\pi} \in (0; 1)$  for each  $\delta f \geq 0$
- monotonically decreasing probability with  $\pi(0) = 1$  and  
 $\lim_{\delta f \rightarrow +\infty} \pi(\delta f) = 0$



Memory can also be used, accepting  $x'$  more easily

if many iterations have elapsed since the last improvement of  $x^*$

# Variable Neighbourhood Search (VNS)

A method very similar to *ILS* is the *Variable Neighbourhood Search* proposed by Hansen and Mladenović (1997)

The main differences between *ILS* and *VNS* are the use of

- the strict acceptance condition:  $f(x') < f(x^*)$
- an **adaptive perturbation mechanism** instead of the fixed one

*VNS* often introduces also neighbourhood modifications

The perturbation mechanism is based on a **hierarchy of neighbourhoods**, that is a **family of neighbourhoods with an increasing parametric size  $k$**

$$N_1 \subset N_2 \subset \dots \subset N_k \subset \dots N_{k_{\max}}$$

Typically one uses the parameterised neighbourhoods

- $N_{H_k}$ , based on the Hamming distance between subsets
- $N_{\mathcal{O}_k}$ , based on the sequences of operations from a basic set  $\mathcal{O}$

and **extracts  $x^{(0)}$  randomly from a neighbourhood of the hierarchy**

# Adaptive perturbation mechanism

It is called *variable neighbourhood* because the neighbourhood used to extract  $x^{(0)}$  varies based on the results of the exchange heuristic

- if a better solution is found, use the smallest neighbourhood, to generate a starting solution very close to  $x^*$  (*intensification*)
- if a worse solution is found, use a slightly larger neighbourhood, to generate a starting solution slightly farther from  $x^*$  (*diversification*)



The method has three parameters

- 1  $k_{\min}$  identifies the smallest neighbourhood to generate new solutions
- 2  $k_{\max}$  identifies the largest neighbourhood to generate new solutions
- 3  $\delta k$  is the increase of  $k$  between two subsequent attempts

The exchange heuristic adopts the smallest neighbourhood to be efficient  
( $N_1$ , or anyway  $N_k$  with  $k \leq k_{\min}$ )

# General scheme of the VNS

*Algorithm* VariableNeighbourhoodSearch( $I, x^{(0)}$ )

$x := \text{SteepestDescent}(x^{(0)}); x^* := x;$

$k := k_{\min};$

*For*  $l := 1$  *to*  $\ell$  *do*

$x' := \text{Shaking}(x^*, k);$

$x' := \text{SteepestDescent}(x');$

*If*  $f(x') < f(x^*)$

*then*  $x^* := x'; k := k_{\min};$

*else*  $k := k + \delta k;$

*If*  $k > k_{\max}$  *then*  $k := k_{\min};$

*EndWhile*;

*Return*  $(x^*, f(x^*));$

- the reference solution  $x^*$  is always the best known solution  $x^*$
- the starting solution is obtained extracting it at random from the current neighbourhood of the reference solution  $N_k(x^*)$
- the exchange heuristic produces a local optimum with respect to the basic neighbourhood  $N$
- if the best known solution improves, the current neighbourhood becomes  $N_{k_{\min}}$
- otherwise, move to a larger neighbourhood  $N_{k+\delta k}$ , never exceeding  $N_{k_{\max}}$ .



# Parameter tuning

The value of  $k_{\min}$  must be

- large enough to get out of the current attraction basin
- small enough to avoid jumping over the adjacent attraction basins

In general, one sets  $k_{\min} = 1$ , and increases it if experimentally profitable

The value of  $k_{\max}$  must be

- large enough to reach any useful attraction basin
- small enough to avoid reaching useless regions of the solution space

Example: the diameter of the search space for the basic neighbourhood:  
 $\min(m, n - m)$  for the *MDP*;  $n$  for the *TSP* and *MAX-SAT*, etc. . .

The value of  $\delta k$  must be

- large enough to reach  $k_{\max}$  in a reasonable time
- small enough to allow each reasonable value of  $k$

In general, one sets  $\delta k = 1$

In order to favour diversification, it is possible to accept  $x'$  when

$$f(x') < f(x^*) + \alpha d_H(x', x^*)$$

where

- $d_H(x', x^*)$  is the Hamming distance fra  $x'$  and  $x^*$
- $\alpha > 0$  is a suitable parameter

This allows to accept worsening solutions as long as they are far away

- $\alpha \approx 0$  tends to accept only improving solutions
- $\alpha \gg 0$  tends to accept any solution

*Of course, the random strategies seen for the ILS can also be adopted*